

# linear\_regrerssion\_corrected\_100%

November 24, 2024

```
[80]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[82]: data = pd.read_csv("boston.csv")
```

```
[84]: data.head()
```

```
[84]:
```

	Unnamed: 0	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	\
0	0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	
1	1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	
2	2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	
3	3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	
4	4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	

	TAX	PTRATIO	B	LSTAT	Price
0	296.0	15.3	396.90	4.98	24.0
1	242.0	17.8	396.90	9.14	21.6
2	242.0	17.8	392.83	4.03	34.7
3	222.0	18.7	394.63	2.94	33.4
4	222.0	18.7	396.90	5.33	36.2

```
[86]: data.describe()
```

```
[86]:
```

	Unnamed: 0	CRIM	ZN	INDUS	CHAS	NOX	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	252.500000	3.613524	11.363636	11.136779	0.069170	0.554695	
std	146.213884	8.601545	23.322453	6.860353	0.253994	0.115878	
min	0.000000	0.006320	0.000000	0.460000	0.000000	0.385000	
25%	126.250000	0.082045	0.000000	5.190000	0.000000	0.449000	
50%	252.500000	0.256510	0.000000	9.690000	0.000000	0.538000	
75%	378.750000	3.677083	12.500000	18.100000	0.000000	0.624000	
max	505.000000	88.976200	100.000000	27.740000	1.000000	0.871000	

	RM	AGE	DIS	RAD	TAX	PTRATIO	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	6.284634	68.574901	3.795043	9.549407	408.237154	18.455534	
std	0.702617	28.148861	2.105710	8.707259	168.537116	2.164946	

min	3.561000	2.900000	1.129600	1.000000	187.000000	12.600000
25%	5.885500	45.025000	2.100175	4.000000	279.000000	17.400000
50%	6.208500	77.500000	3.207450	5.000000	330.000000	19.050000
75%	6.623500	94.075000	5.188425	24.000000	666.000000	20.200000
max	8.780000	100.000000	12.126500	24.000000	711.000000	22.000000

	B	LSTAT	Price
count	506.000000	506.000000	506.000000
mean	356.674032	12.653063	22.532806
std	91.294864	7.141062	9.197104
min	0.320000	1.730000	5.000000
25%	375.377500	6.950000	17.025000
50%	391.440000	11.360000	21.200000
75%	396.225000	16.955000	25.000000
max	396.900000	37.970000	50.000000

```
[88]: data.isnull().sum()
```

```
[88]: Unnamed: 0      0
      CRIM          0
      ZN           0
      INDUS        0
      CHAS         0
      NOX          0
      RM           0
      AGE          0
      DIS          0
      RAD          0
      TAX          0
      PTRATIO      0
      B            0
      LSTAT        0
      Price        0
      dtype: int64
```

```
[90]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0  506 non-null    int64
1   CRIM        506 non-null    float64
2   ZN          506 non-null    float64
3   INDUS       506 non-null    float64
4   CHAS        506 non-null    float64
5   NOX         506 non-null    float64
```

```

6   RM          506 non-null   float64
7   AGE         506 non-null   float64
8   DIS         506 non-null   float64
9   RAD         506 non-null   float64
10  TAX         506 non-null   float64
11  PTRATIO     506 non-null   float64
12  B           506 non-null   float64
13  LSTAT       506 non-null   float64
14  Price       506 non-null   float64
dtypes: float64(14), int64(1)
memory usage: 59.4 KB

```

```
[92]: data.drop('Unnamed: 0', axis = 1,inplace=True)
```

```
[94]: data.head()
```

```
[94]:
```


	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT	Price
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

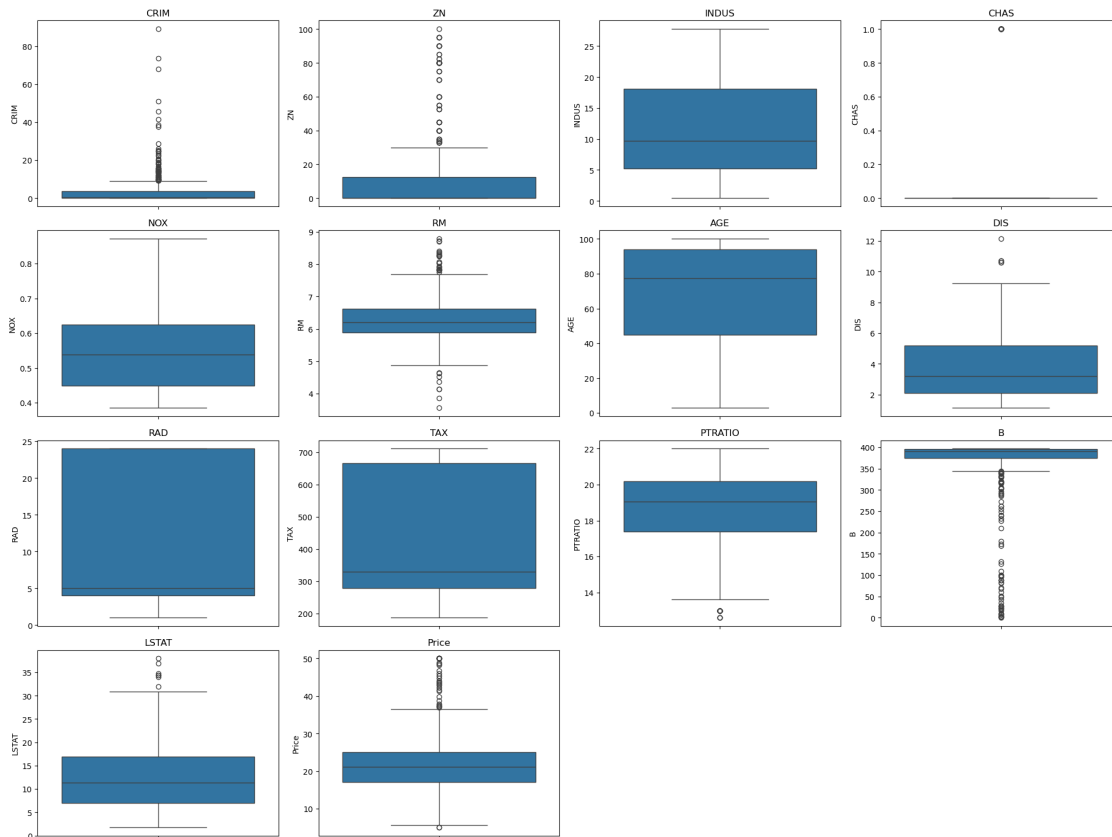
```
[99]: sns.boxplot(x='LSTAT', data=data)
```

```
[99]: <Axes: xlabel='LSTAT'>
```



```
[105]: plt.figure(figsize=(20, 15)) # Adjust figure size as needed
for i, column in enumerate(data.columns, 1):
    plt.subplot(4, 4, i) # Adjust grid dimensions based on the number of
    ↪ columns
    sns.boxplot(y=data[column])
    plt.title(column)

plt.tight_layout()
plt.show()
```



```
[107]: import pandas as pd

# Function to remove outliers using the IQR method
def remove_outliers(df):
    numeric_columns = df.select_dtypes(include='number').columns
    cleaned_df = df.copy()
    for column in numeric_columns:
        Q1 = df[column].quantile(0.25) # First quartile (25th percentile)
        Q3 = df[column].quantile(0.75) # Third quartile (75th percentile)
        IQR = Q3 - Q1 # Interquartile range
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        # Filter the dataset to include only data within bounds
        cleaned_df = cleaned_df[(cleaned_df[column] >= lower_bound) &
        ↪ (cleaned_df[column] <= upper_bound)]
    return cleaned_df

# Apply the function to the dataset
data = remove_outliers(data)
```

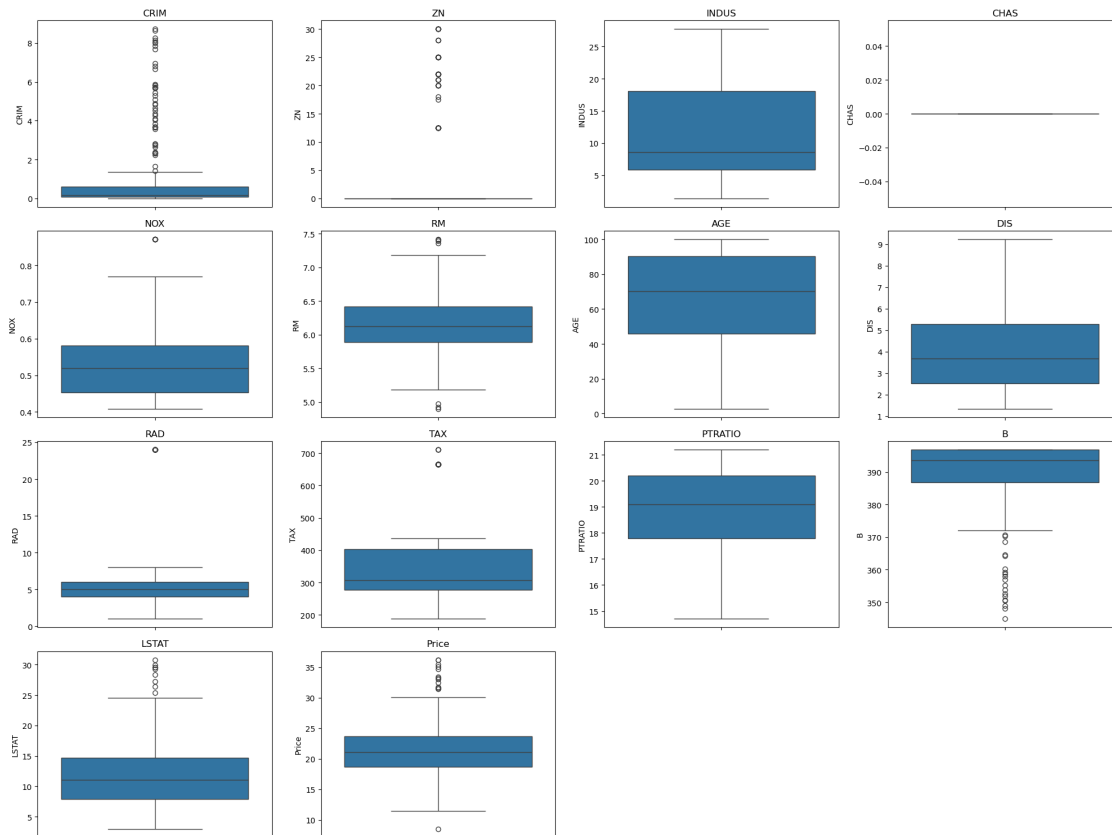
```
# Check the resulting dataset
print("Original dataset shape:", data.shape)
print("Cleaned dataset shape:", data.shape)
```

Original dataset shape: (268, 14)

Cleaned dataset shape: (268, 14)

```
[109]: plt.figure(figsize=(20, 15)) # Adjust figure size as needed
for i, column in enumerate(data.columns, 1):
    plt.subplot(4, 4, i) # Adjust grid dimensions based on the number of
    ↪ columns
    sns.boxplot(y=data[column])
    plt.title(column)

plt.tight_layout()
plt.show()
```



```
[115]: # Calculate the IQR for the 'AGE' column
Q1 = data['AGE'].quantile(0.25) # 25th percentile
Q3 = data['AGE'].quantile(0.75) # 75th percentile
IQR = Q3 - Q1
```

```

# Define lower and upper bounds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Replace outliers with 0 in the 'AGE' column
data['AGE'] = data['AGE'].apply(lambda x: 0 if x < lower_bound or x > upper_bound else x)

# Check the updated column
data['AGE'].head()

```

```

[115]: 0    65.2
      1    78.9
      2    61.1
      3    45.8
      4    54.2
      Name: AGE, dtype: float64

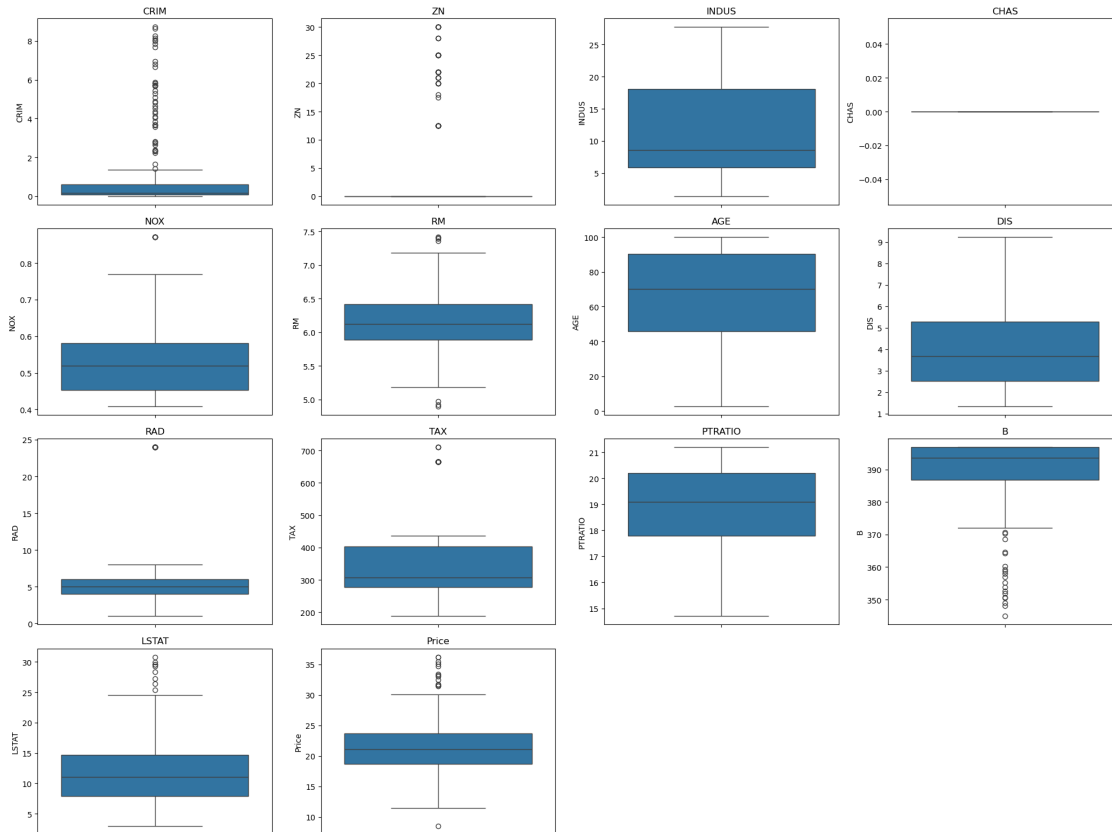
```

```

[117]: plt.figure(figsize=(20, 15)) # Adjust figure size as needed
      for i, column in enumerate(data.columns, 1):
          plt.subplot(4, 4, i) # Adjust grid dimensions based on the number of
          ↪ columns
          sns.boxplot(y=data[column])
          plt.title(column)

      plt.tight_layout()
      plt.show()

```



```
[123]: X = data[['AGE', 'RM', 'LSTAT']]
       y = data['Price']
```

```
[125]: # Calculate the IQR for the 'AGE' column
Q1 = data['RM'].quantile(0.25) # 25th percentile
Q3 = data['RM'].quantile(0.75) # 75th percentile
IQR = Q3 - Q1

# Define lower and upper bounds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Replace outliers with 0 in the 'AGE' column
data['RM'] = data['RM'].apply(lambda x: 0 if x < lower_bound or x > upper_bound
                               else x)

# Check the updated column
data['RM'].head()
```



```
[125]: 0    6.575
      1    6.421
      2    7.185
      3    6.998
      4    7.147
      Name: RM, dtype: float64
```

```
[127]: # Calculate the IQR for the 'AGE' column
Q1 = data['LSTAT'].quantile(0.25) # 25th percentile
Q3 = data['LSTAT'].quantile(0.75) # 75th percentile
IQR = Q3 - Q1

# Define lower and upper bounds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

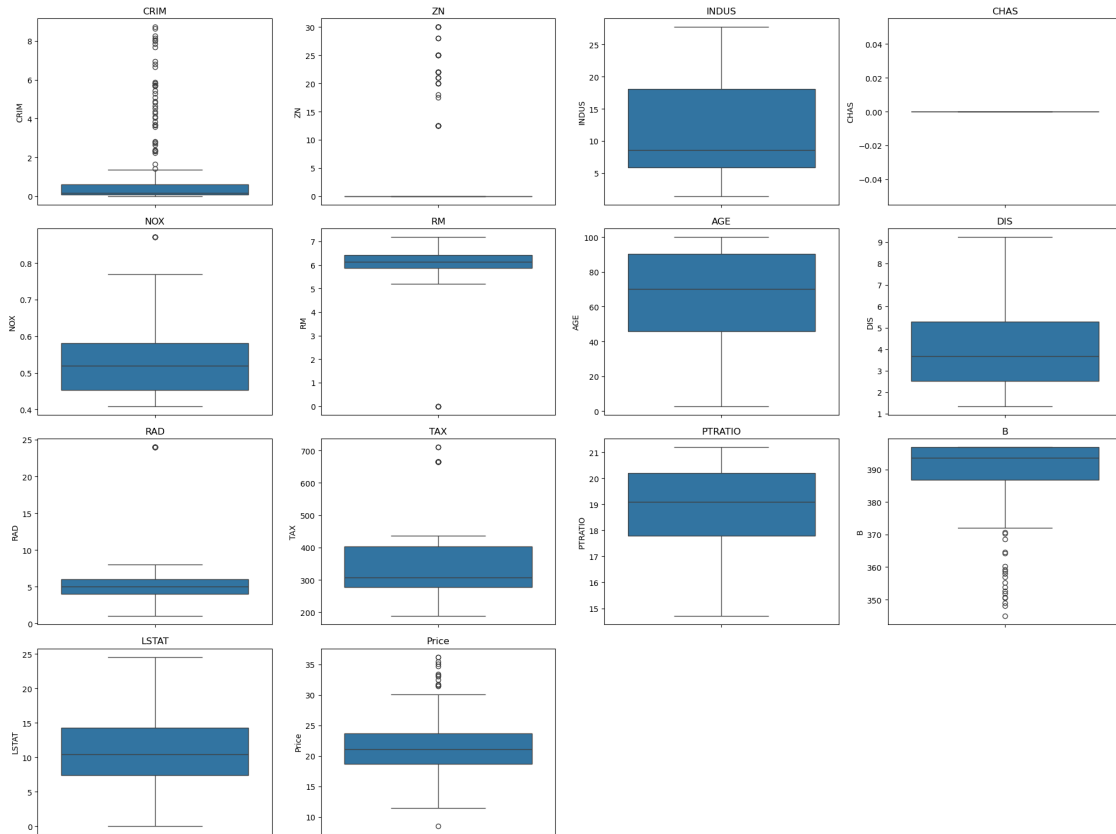
# Replace outliers with 0 in the 'AGE' column
data['LSTAT'] = data['LSTAT'].apply(lambda x: 0 if x < lower_bound or x > upper_bound else x)

# Check the updated column
data['RM'].head()
```

```
[127]: 0    6.575
      1    6.421
      2    7.185
      3    6.998
      4    7.147
      Name: RM, dtype: float64
```

```
[129]: plt.figure(figsize=(20, 15)) # Adjust figure size as needed
for i, column in enumerate(data.columns, 1):
    plt.subplot(4, 4, i) # Adjust grid dimensions based on the number of columns
    sns.boxplot(y=data[column])
    plt.title(column)

plt.tight_layout()
plt.show()
```



```
[131]: X = data[['AGE', 'RM', 'LSTAT']]
       y = data['Price']
```

```
[133]: from sklearn.model_selection import train_test_split
```

```
[135]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
       ↪ random_state=40)
```

```
[139]: X_train.shape
```

```
[139]: (214, 3)
```

```
[141]: X_test.shape
```

```
[141]: (54, 3)
```

```
[143]: y_train.shape
```

```
[143]: (214,)
```

```
[145]: y_test.shape
```

[145]: (54,)

```
[149]: from sklearn.linear_model import LinearRegression  
       from sklearn.metrics import r2_score, mean_squared_error
```

```
[151]: model = LinearRegression()
```

```
[153]: model.fit(X_train, y_train)
```

[153]: LinearRegression()

```
[158]: y_pred = model.predict(X_test)
```

```
[160]: mse = mean_squared_error(y_test, y_pred)  
       r2 = r2_score(y_test, y_pred)  
  
       print(f"Mean Squared Error: {mse}")  
       print(f"R-squared Score: {r2}")
```

Mean Squared Error: 12.2986764374419

R-squared Score: 0.4727964006739217

```
[ ]:
```